

VISVESVARAYA TECHNOLOGICAL UNIVERSITY
JNANASANGAMA, BELAGAVI - 590018



A Minor Project Report On
VibeCheck Sentiment Intelligence

Submitted in partial fulfillment for the AAT in MLOps (22AI61)
In
Artificial Intelligence and Machine Learning Department

Submitted by

USMANAGANI DAVALASAB ATTAR (1DS23AI062)

VISHUNU R (1DS23AI063)

KEERTI T NAIK (1DS24AI400)

MANJUNATH (1DS24AI401)

Guide: Mr. Basudev Panda
Course Coordinator: Prof. Anupama V P



DAYANANDA SAGAR COLLEGE OF ENGINEERING

(An Autonomous Institute Affiliated to VTU, Belagavi.) Shavige

Malleshwara Hills, Kumaraswamy Layout, Bengaluru-560111

Department of Artificial Intelligence and Machine Learning
2025 – 2026

ACKNOWLEDGEMENT

The success and outcome of this project require the guidance and assistance of many people. We would like to add a few words of appreciation for the people who have been part of this project right from its inception, without their support patience and guidance the task would not have been completed. It is to them we owe them our deepest gratitude.

We are grateful to **Dr. B G Prasad**, Principal, Dayananda Sagar College of Engineering, for providing an opportunity to do this project as a part of our curriculum and for his kind cooperation for the project.

We are very much grateful to **Dr. Vindhya P Malagi**, Professor and HOD, Department of Artificial Intelligence and Machine Learning, Dayananda Sagar College of Engineering, Bangalore for providing the encouragement for completion of our project.

We would like to express our deep gratitude to our Project Guide, **Mr. Basudev Panda**. For his valuable guidance, patience, constant supervision and timely suggestions provided in making of this project and for his support throughout the completion.

We are also thankful to our parents and friends for their constant help and constructive suggestions throughout our project.

USMANAGANI DAVALASAB ATTAR (1DS23AI062)

VISHUNU R (1DS23AI063)

KEERTI T NAIK (1DS24AI400)

MANJUNATH (1DS24AI401)

ABSTRACT

In the era of big data, the ability to systematically interpret human emotion from unstructured textual data has become a critical requirement for organizations seeking real-time consumer insights. This project details the design, implementation, and deployment of a robust Machine Learning Operations (MLOps) pipeline for automated Sentiment Analysis. Utilizing a comprehensive dataset featuring over 30 granular micro-emotions, we engineered a preprocessing architecture that aggregates these into three macro-sentiment categories: Positive, Negative, and Neutral. The pipeline utilizes Term Frequency-Inverse Document Frequency (TF-IDF) vectorization with unigram and bigram configurations to capture semantic nuances across 6,000 feature dimensions. A comparative empirical study was conducted across five distinct machine learning paradigms: Logistic Regression, Multinomial Naive Bayes, Linear Support Vector Classification (SVC), Random Forest, and Decision Trees. Evaluation metrics—including Precision, Recall, and F1-Score—revealed that linear methodologies, specifically Linear SVC, significantly outperformed non-linear tree-based models, which struggled with the high-dimensional sparsity of the textual feature space. To further enhance predictive stability and mitigate individual algorithmic biases, a Voting Ensemble architecture was developed, aggregating the strengths of the top-performing classifiers. The finalized system was serialized and deployed via a Streamlit-based graphical user interface (GUI), providing a seamless, real-time inference environment for end-users. This project demonstrates that the integration of rigorous NLP preprocessing with synchronized ensemble learning offers a scalable and reliable solution for linguistic intelligence extraction.

TABLE OF CONTENTS

ACKNOWLEDGEMENT	i
ABSTRACT	ii
1 INTRODUCTION	
1.1 Overview of Mlops	01
1.2 Problem Statement	02
1.3 Objectives	02
1.4 Methodology	03
2 MODELS SELECTION	
2.1 Model 1	04
2.2 Model 2	05
2.3 Model 3	05
2.4 Model 4	06
2.5 Model 5	06
3 IMPLEMENTATION	
3.1 Pipeline Design	07
3.1.1 Data Collection	
3.1.2 Data Analysis	
3.1.3 Training And Testing The Model	
3.1.4 Deployment Of The Model	
3.2 Performance comparision of Models Used	09
4 SDG MAPPED	10
5 RESULTS	11
CONCLUSION	14

1 INTRODUCTION

1.1 Overview Of MLOps

Machine Learning Operations (MLOps) represents a fundamental paradigm shift in how artificial intelligence systems are developed, deployed, and maintained. Historically, machine learning models were built in isolated, static environments by data scientists. While these models might achieve high accuracy during the experimental phase, transitioning them into production environments where they must process continuous, live data often resulted in systems that were fragile, unscalable, or prone to degradation over time. MLOps emerged as a critical engineering discipline to bridge the gap between model development and software operations.

At its core, MLOps is an extension of DevOps tailored specifically for machine learning pipelines. It incorporates the principles of Continuous Integration (CI), Continuous Deployment (CD), and Continuous Training (CT). Unlike traditional software development, where the code dictates the behavior of the application, machine learning applications are governed by both code and data. As real-world data constantly evolves—a phenomenon known as "data drift" or "concept drift"—models inevitably lose their predictive power over time. MLOps establishes automated pipelines that monitor model performance, trigger alerts when accuracy thresholds are breached, and facilitate rapid retraining and redeployment of models.

In the context of this project, implementing MLOps principles ensures that our sentiment analysis tool is not merely a static python script, but a robust pipeline. From data ingestion and preprocessing to model serialization and cloud deployment via Streamlit, the MLOps framework guarantees that the application can receive new textual data, update its internal parameters, and serve reliable predictions to end-users without service interruption. It encapsulates best practices in version control (tracking both code and model assets), automated testing, and scalable serving.

1.2 Problem Statement

In the modern digital era, human interaction, feedback, and opinions are overwhelmingly expressed through unstructured textual data on social media platforms, product review forums, and customer support channels. For organizations, extracting actionable insights from this colossal volume of unstructured data is computationally daunting. Human language is inherently complex, characterized by sarcasm, slang, domain-specific terminology, and highly nuanced emotional variations.

The primary problem this project addresses is the automated interpretation of subjective text sequences. Specifically, how do we systematically map an incredibly diverse array of human emotions—ranging from "euphoria" and "gratitude" to "rage" and "melancholy" into actionable, broad sentiment categories? Without a highly scalable automated system, organizations cannot process user feedback in real-time, leading to slow response times and a poor understanding of public perception.

This project aims to solve this by engineering a comprehensive Sentiment Analysis pipeline that classifies arbitrary user text into 'Positive', 'Negative', or 'Neutral' buckets. The challenge lies in managing the high-dimensional, sparse nature of textual data, mitigating dataset imbalances, and delivering a reliable prediction engine. The deployed system must overcome linguistic noise and remain easily accessible for end-users, requiring a combination of robust Natural Language Processing (NLP) techniques and optimized machine learning algorithms.

1.3 Objectives

The development and deployment of this MLOps sentiment analysis pipeline are driven by the following core objectives:

1. Automated Data Processing: To engineer a robust ingestion mechanism capable of cleaning, normalizing, and standardizing raw, unstructured text data. This includes mapping a complex array of highly specific emotion labels into three fundamental macro-categories: Positive, Negative, and Neutral.
2. Advanced Feature Engineering: To implement mathematically rigorous text vectorization techniques, specifically Term Frequency-Inverse Document Frequency (TF-IDF), utilizing a combination of unigrams and bigrams to capture contextual semantics effectively.
3. Comprehensive Model Evaluation: To train, evaluate, and rigorously compare at least five distinct machine learning classification algorithms. This comparative study aims to understand how linear, probabilistic, and tree-based models respond to high-dimensional text data.
4. Enhanced Predictive Stability: To mitigate the inherent biases and variance of individual algorithms by engineering a Voting Ensemble classifier that aggregates the predictive power of the top-performing models.

1.4 Methodology

To achieve the project objectives, a structured, widely accepted machine learning lifecycle methodology was adapted. This methodology ensures traceability, robustness, and reproducibility of the results. The workflow is divided into several iterative phases:

Phase 1: Data Ingestion and Transformation The initial phase involves acquiring the `sentimentdataset.csv`. The raw dataset contains highly granular, specific emotional tags. An automated mapping function was developed to assign each specific emotion explicitly to broader 'Positive', 'Negative', or 'Neutral' buckets. All textual data was subjected to preliminary cleaning, which included stripping leading/trailing whitespace, standardizing string encodings, and converting the target variable into a categorical structure suitable for supervised learning algorithms.

Phase 2: Text Vectorization Strategy Machine learning algorithms cannot naturally interpret raw strings. We utilized `TfidfVectorizer` to quantify the text. The vectorization was configured to cap at a substantial 6,000 maximum features to prevent memory exhaustion while retaining crucial vocabulary. We implemented an `ngram_range` of (1, 2) to capture bigrams, allowing the models to understand sequential word pairs rather than isolating every word out of context. The `sublinear_tf=True` parameter was utilized to logarithmically scale term frequencies, which has empirically been shown to improve accuracy in varying-length texts.

Phase 3: Model Training and Hyperparameter Configuration An 80-20 Train/Test split algorithm was utilized (`train_test_split`), heavily leveraging the `stratify` parameter to guarantee that our training and testing data splits contained a representative proportion of the overarching target classes. Five distinct baseline classifiers were established. To combat the natural class imbalance found in the sentiment dataset, the `class_weight` is equals to 'balanced' heuristic was enabled across models. This mathematically forces the algorithm to pay more attention (give a higher penalty for misclassification) to minority classes.

Phase 4: Evaluation and Ensemble Engineering Post-training, models were subjected to evaluation using unseen data from the test split. We extracted precision, recall, F1-scores, and overall accuracy. Confusion matrices were generated to observe exact misclassification behaviors. Based on the evaluation, an Ensemble Voting Classifier was constructed.

Phase 5: Model Serialization and Deployment The finalized vectorizer geometry and the trained Voting Ensemble model state were serialized into lightweight binary formats (.pkl files) using Python's `joblib`. These assets were then integrated into a Streamlit web application (`app.py`), resulting in an architecture that separates the heavy computational training processes from the lightweight, high-speed inference constraints required by the end-user dashboard.

2 MODEL SELECTION

Reasons for particular model selection

The process of model selection in Natural Language Processing is guided by the fundamental "No Free Lunch" theorem of machine learning, which posits that no single algorithm works best for every problem. Sentiment analysis problems are uniquely challenging because the vectorization of language produces high-dimensional, sparse matrices. Rather than choosing an algorithm blindly, a diverse portfolio of five structurally different models was deliberately selected.

We included linear models (Logistic Regression, Linear SVC) because they are historically exceptional at handling massive numbers of sparse features without overfitting. A probabilistic model (Naive Bayes) was selected as a traditional computationally lightweight NLP baseline to measure against.

Finally, non-linear tree-based models (Random Forest, Decision Tree) were introduced to identify if complex conditional interactions between specific words were dictating sentiment more strongly than simple linear combinations. This comprehensive coverage ensures that the eventual production model is chosen based on empirical rigor rather than theoretical assumptions.

2.1 Logistic Regression

Logistic Regression is a foundational classification algorithm from the generalized linear model family. Despite its name, it is utilized exclusively for classification. It works by passing a linear combination of input features through a Sigmoid (logistic) function, which squashes the output into a probability score bounded strictly between 0 and 1.

Advantages:

- **Interpretability:** The coefficients assigned to the TF-IDF features directly represent importance, making it incredibly clear which words drive positive or negative classification.
- **Probabilistic Calibration:** Because it outputs raw probability scores, we can confidently set certainty thresholds.
- **Efficiency:** Logistic Regression is famously computationally light to train, allowing rapid iteration on the vast 6,000-dimensional matrix created by our text vectorizer.

Disadvantages:

- **Linear Boundaries:** The core mathematical assumption of Logistic Regression is that a straight line (or hyperplane in higher dimensions) can cleanly separate the classes. If sentiment rules are highly non-linear or rely heavily on complex word interdependencies, Logistic Regression will underperform.

2.2 Naive Bayes (MultinomialNB)

The Multinomial Naive Bayes algorithm relies on the principles of Bayesian probability. It calculates the conditional probability of each class given the observed words in the document. The term "multinomial" refers to its optimization for discrete counts or fractional counts (like TF-IDF scores). The algorithm relies on maximum a posteriori (MAP) estimation to deduce sentiment.

Advantages:

- **Simplicity and Speed:** Naive Bayes requires only a single pass over the training data to calculate probabilities, making it the fastest algorithm to train.
- **Performance with Small Datasets:** It naturally converges quickly, often outperforming complex algorithms when training data is highly limited.
- **Robustness to Irrelevant Features:** Words that do not correlate strongly with sentiment generally wash out in the probability calculations without negatively impacting the model.

Disadvantages:

- **The Naive Assumption:** It naively assumes that every feature (word or bigram) is completely independent of every other feature. In human language, this is false—words string together to create heavily dependent contexts. This assumption can sometimes limit its ceiling regarding absolute accuracy, especially on long, complex sentences.

2.3 Linear SVC (Support Vector Classifier)

Linear Support Vector Classification is a mathematical algorithm that seeks to find the "optimal hyperplane" a decision boundary in N dimensional space that maximizes the geometric distance (margin) between the outermost data points (support vectors) of different classes. The linear kernel specifically utilizes dot products without projecting the data into higher dimensional planes.

Advantages:

- **High Dimensional Superiority:** Linear SVC is arguably the strongest traditional ML algorithm for text classification because it inherently mitigates the curse of dimensionality. Support vectors define the boundary, meaning irrelevant features have zero impact on the final model.
- **Robust to Overfitting:** The regularization parameter (C) acts as a strict penalty on absolute boundary fitting, ensuring the model generalizes incredibly well on unseen evaluation text.

Disadvantages:

- **No Native Probabilities:** Unlike Logistic Regression, SVC outputs point-in-space distances rather than probabilities. Obtaining strict confidence scores requires computationally expensive Platt scaling.
- **Sensitivity to Outliers:** Substantial text anomalies mapped near the margin can sometimes dramatically skew the hyperplane calculation.

2.4 Random Forest

Random Forest is an ensemble learning method constructed upon the foundations of Decision Trees. It mitigates the flaws of individual trees by utilizing "bagging" (Bootstrap Aggregating). It builds hundreds of deep decision trees using random subsets of both the training rows and the TF-IDF feature columns. During inference, the forest takes a democratic vote among all constituent trees to finalize a prediction.

Advantages:

- **Non-Linear Capability:** Random Forest effortlessly maps incredibly complex, non-linear boundaries. If a specific sentiment only triggers when word A exists AND word B does not exist, the tree logic perfectly captures this.
- **Variance Reduction:** Because individual trees are decorrelated by utilizing variable feature subsets, the massive variance issues standard to decision trees are effectively neutralized.

Disadvantages:

- **Sparsity Struggles:** Tree-based models generally struggle on sparse, high-dimensional arrays heavily populated by zeros (like TF-IDF vectors). Splitting logic becomes highly inefficient.
- **Computational Overhead:** Storing and traversing hundreds of trees consumes substantial memory and significantly slows down inference speeds during deployment compared to linear models.

2.5 Decision Tree


A Decision Tree is a non-parametric model that systematically splits data into partitions by evaluating information gain or Gini impurity at specific feature threshold levels. It results in a highly readable flowchart-like structure, asking binary questions about word occurrences traversing from root down to leaf nodes consisting of class labels.

Advantages & Disadvantages:

- **Total Transparency:** The logical path taken to reach a prediction is entirely visible. We can precisely trace which words caused the tree to fork to a negative sentiment.
- **Extreme Overfitting Risk:** If depth is not heavily constrained, a Decision Tree will continue splitting until it has essentially memorized the training text data verbatim. This results in catastrophic failure on new, unseen test data.

3 IMPLEMENTATION

3.1 Pipeline Design (Diagram To Be Added)



The pipeline design for this MLOps architecture has been carefully structured to isolate data preprocessing, complex model training logic, and user-facing web deployment. The design is modular, meaning the data pipeline can be modified without needing to rewrite the dashboard frontend, maintaining standard engineering loosely coupled constraints.

3.1.1 Data Collection

The foundation of effective machine learning is a robust dataset. The dataset utilized for this project (sentimentdataset.csv) features a comprehensive arrangement of real-world text entries derived from various conversational and informal mediums. It encompasses thousands of rows containing unstructured text and corresponding emotional classifications. The source data was particularly characterized by an expansive variety of micro-emotions (e.g., ecstasy, terror, grief, bliss). Due to the inherent difficulty in building reliable statistical boundaries across 30+ separate emotions, the primary MLOps preprocessing script was engineered to aggregate these specific tags via a sweeping dictionary mapping structure strictly into a simplified macro-classification scheme: Positive, Negative, or Neutral.

3.1.2 Data Analysis

Once the data labels were mapped into buckets, rigorous analytical processes were applied. Preprocessing removed unnecessary noise such as special characters, HTML fragments, or erratic punctuation which offer minimal predictive value and inflate the dimensionality of the vector space. Exploratory Data Analysis (EDA) revealed typical class distributions, where neutral or broadly positive textual phrases frequently outnumbered deeply negative statements.

This identified class imbalance mandated the use of stratify=y during data splitting procedures, ensuring that the critical but rare negative comments were evenly distributed into the validation test suite rather than being clustered heavily in the training batch, which would generate misleading evaluation metrics.

3.1.3 Training And Testing The Model

The core implementation procedure involved transforming the cleaned text into a format actionable by algorithms. The TF-IDF methodology was selected over simpler bag-of-words counting because it actively penalizes extremely common English stop words ("the", "and", "is") while exponentially weighting specific, rare words ("terrible", "magnificent") that carry heavy sentiment payload. The data was split into an 80% training suite to teach the math, reserving 20% strictly for unbiased testing.

To comprehensively evaluate performance, we moved past mere Accuracy which can be deeply misleading on imbalanced datasets. We utilized Precision (reducing false positives) and Recall (reducing false negatives) to calculate the F1-Score, providing the harmonic mean that signifies true robust generalization.

3.1.4 Deployment Of The Model

An algorithm is functionally useless if analysts cannot access its intelligence. Deployment was realized via Streamlit, a robust application framework tailored for deploying Python-centric data operations. Through a custom app.py script, an aesthetic User Interface (UI) was constructed utilizing custom CSS for modern glassmorphism design. The script deserializes the optimal model (best_sentiment_model.pkl) and the pre-fitted text geometry (tfidf_vectorizer.pkl) directly into active memory. When an end-user inputs custom text into the dashboard, it is dynamically sanitized, passed through vectorization mathematics, and pushed through the Voting Ensemble to render instantaneous sentiment classification on the screen with corresponding confidence probabilities.

The application loads the serialized machine learning model (best_sentiment_model.pkl) along with the pre-trained TF-IDF vectorizer (tfidf_vectorizer.pkl) directly into memory for efficient real-time predictions. When a user enters text into the dashboard, the input is automatically cleaned, transformed into numerical vectors using TF-IDF, and processed through the Voting Ensemble classifier. The system then instantly displays the predicted sentiment along with confidence scores, enabling fast and accurate sentiment analysis through an accessible interface.

3.2 Performance Comparison Of Models Used

The empirical results of the test split evaluations across the five distinct classifiers are detailed in the comprehensive metrics table below. These metrics strictly represent performance on the 20% validation chunk that the models had never seen during active training parameters optimization.

Models	Precision	Recall	F1 score	Accuracy
Model 1: Logistic Regression	0.7291	0.7211	0.7192	0.7211
Model 2: Naive Bayes	0.7020	0.6871	0.6572	0.6871
Model 3: Linear SVC	0.7415	0.7347	0.7331	0.7347
Model 4: Random Forest	0.7500	0.6599	0.6324	0.6599
Model 5: Decision Tree	0.5579	0.4898	0.3692	0.4898

Analysis of Performance Comparison: Analyzing the performance matrix clearly demonstrates the profound superiority of linear algorithms dealing specifically with TF-IDF configurations. Model 3 (Linear SVC) decisively outperformed all standalone models, attaining the highest unified F1 Score of 0.7331. Model 1 (Logistic Regression) remained exceptionally competitive with a reliable 0.7192 F1 Score, highlighting its fundamental strength in processing semantic weight linearly. By brutal contrast, Model 5 (Decision Tree) suffered near catastrophic failure at generalizing the complex vocabulary, resulting in severely degraded recall and precision metrics that hover slightly below 0.50—indicating that static hierarchy thresholds cannot accurately encapsulate the massive variability of human colloquial phrasing.

The evaluation results indicate that models based on linear classification techniques perform significantly better when combined with TF-IDF feature representations. Since TF-IDF generates high-dimensional sparse vectors, algorithms such as Logistic Regression and Linear SVC are more effective at identifying meaningful decision boundaries within textual data. Among all tested classifiers, Linear SVC achieved the strongest overall balance between precision, recall, and F1 score, demonstrating its capability to accurately capture contextual sentiment patterns while minimizing classification errors. Logistic Regression also delivered highly stable and competitive performance, reinforcing its reliability for sentiment-based natural language processing tasks.

In contrast, tree-based approaches such as Decision Tree and Random Forest showed comparatively weaker generalization capabilities on the dataset. These models struggled to efficiently interpret the complex semantic relationships and vocabulary diversity present in human language, resulting in lower recall and inconsistent prediction accuracy. The poor performance of the standalone Decision Tree model particularly highlights the limitations of rigid hierarchical splitting methods when applied to unstructured textual data.

4 SDG MAPPED

This Sentiment Intelligence project leverages advanced Natural Language Processing (NLP) and Machine Learning to process and analyze textual data. The development and potential applications of this system align with the following UN Sustainable Development Goals

SDG 3: Good Health and Well-being:

The project contributes significantly to SDG 3 by supporting mental health awareness and digital well-being through intelligent sentiment monitoring. In today's digital era, social media platforms and online communication channels often reflect the emotional state of individuals and communities. By accurately analyzing emotional tone, sentiment polarity, and subjectivity in real time, the sentiment analyzer can act as an early-warning system for identifying signs of emotional distress, anxiety, frustration, or negative behavioral trends.

This capability has strong applications in the field of public health and psychological well-being. Governments, healthcare organizations, and researchers can use the system to monitor large-scale emotional trends across digital platforms and identify sudden increases in negative sentiment associated with stress, depression, or crisis situations. Such analysis can help institutions respond proactively by launching awareness campaigns, providing mental health support, or investigating the causes of widespread distress.

SDG 9: Industry, Innovation, and Infrastructure:


This project strongly aligns with SDG 9 by showcasing the practical application of advanced technologies to build an intelligent and scalable digital solution. The sentiment analysis system is not limited to a basic machine learning implementation; instead, it incorporates a hybrid AI architecture that combines a Voting Ensemble model with heuristic Natural Language Processing (NLP) techniques. This approach enhances prediction accuracy, robustness, and adaptability by leveraging the strengths of multiple algorithms rather than relying on a single classifier. Such innovation reflects the growing trend in modern AI systems toward ensemble intelligence and intelligent automation.

SDG 16: Peace, Justice, and Strong Institutions:


This project supports SDG 16 by promoting transparency, inclusivity, and safer digital communication environments through sentiment intelligence. In democratic societies and digital communities, understanding public opinion is essential for ensuring trust, accountability, and effective communication between institutions and citizens. The sentiment analysis system provides a data-driven mechanism for analyzing how people react to policies, events, campaigns, or public discussions, thereby enabling institutions to make more informed and responsive decisions.

5 RESULTS


INDIVIDUAL MODEL PERFORMANCE GRAPHS: CONFUSION MATRICES




Logistic Regression Confusion Matrix: This matrix visually illustrates the classification distribution. We can observe high values down the central diagonal, representing true accuracy. The heat-mapping colors denote concentration densities, showcasing Logistic Regression's strong capability at distinctly isolating Positive from Negative contexts with limited mid-range spillage.




Multinomial Naive Bayes Confusion Matrix: Here we can visualize the impact of the algorithm's probabilistic feature-independence assumptions. While it establishes a reasonable baseline, the off-diagonal cells reveal a slightly elevated rate of false classifications compared to its linear constraints counterparts, especially concerning complex ambiguous phrases. This behavior occurs because the Naive Bayes algorithm assumes that all textual features are conditionally independent of one another, an assumption that rarely holds true in real-world natural language processing tasks.



Linear SVC Confusion Matrix: This visualization verifies Linear SVC as the supreme individual algorithm. The exceptionally dark blue concentrations spanning strictly upon the True Positive matching diagonal explicitly demonstrate the successful execution of its large-margin optimization parameters over the sparse textual mapping. The confusion matrix further highlights the model's ability to maintain balanced classification performance across multiple sentiment categories with minimal misclassification leakage into adjacent classes.




Random Forest Confusion Matrix: While the Random Forest effectively managed to drastically reduce the variance issues of single trees, the matrix highlights its intrinsic struggles with precision-recall trade-offs when navigating vast arrays of high-dimensional zeros typical in our sparse textual vector configurations.



Decision Tree Confusion Matrix: This visualization starkly represents severe overfitting consequences. The extensive dispersion of values outside the central diagonal confirms that the strict Boolean mapping failed disastrously on text it had not explicitly memorized during training, generating extensive misclassifications.

COLLAGE OF ALL MODELS PERFORMANCE GRAPHS



Aggregate Performance Bar Chart: This comprehensive collage graph summarizes the overarching investigation. It visually aligns the F1 scores alongside total Accuracy and Precision metrics across all algorithms. The disparities are readily apparent, conclusively verifying the structural advantages inherent in linear classification methodologies concerning Natural Language Processing operations.

The visualization clearly demonstrates the consistent dominance of Linear SVC and Logistic Regression across nearly all evaluation metrics, emphasizing their reliability and stability in sentiment classification tasks. In contrast, non-linear and tree-based models exhibit noticeable performance fluctuations, indicating weaker adaptability to sparse and high-dimensional textual representations generated through TF-IDF vectorization.

CONCLUSION

The empirical investigations administered throughout the course of this MLOps pipeline lifecycle rendered conclusive insights regarding algorithm applicability to semantic intelligence extraction. The rigorous comparison demonstrated that not all machine learning paradigms handle vast NLP matrices equitably. Linear SVC (Support Vector Classifier) proved to be the superior predictive mechanism within this study context, achieving strong generalized performance across Precision, Recall, and Accuracy.

The Architecture Advantage: The prevailing reason Linear SVC excelled lies within the architectural properties of TF-IDF vectors. The vectorization phase transformed raw phrasing into continuous 6,000-dimensional arrays populated mostly by zero values, since no single sentence contains all 6,000 words. Non-linear algorithms like Decision Tree and Random Forest often degrade when splitting data across thousands of sparse features, which can lead to overfitting on the training vocabulary. Conversely, Linear SVC navigates the curse of dimensionality effectively. It functions by calculating an optimal dividing hyperplane, maximizing geometric separation while focusing on the most relevant sentimental tokens.

The Solution – Engineering an Ensemble: However, despite identifying a prime standalone algorithm, advanced Machine Learning Operations recognizes that individual models possess structural biases. Although Linear SVC attained the highest validation scores, relying on a singular hypothesis boundary may create fragility in production environments.

Therefore, our concluding operational decision was to deploy a Voting Ensemble architecture. Recognizing the strengths of Linear SVC (margin maximization), Logistic Regression (calibrated linear approximations), and Random Forest (complex variable relationships), the ensemble unites these architectures collaboratively. Incoming text is evaluated by all three algorithms simultaneously, and the final deployed configuration functions under a majority-rules voting protocol. By combining distinct paradigms, the Voting Ensemble neutralizes many blind spots of individual classifiers while supporting a robust real-time Sentiment Intelligence Dashboard on Streamlit.

The ensemble-based deployment strategy also improves prediction stability when processing unseen or linguistically complex user inputs in real-world environments. By aggregating outputs from multiple classifiers, the system reduces dependency on a single decision boundary and minimizes the risk of biased or inconsistent sentiment predictions. This collaborative architecture enhances robustness, scalability, and fault tolerance, making the model more reliable for continuous real-time inference tasks. Ultimately, the integration of ensemble intelligence with MLOps deployment practices establishes a production-ready framework capable of delivering accurate and sustainable sentiment analysis at scale.